



3dssd — single-model baseline

MPS off

MPS on

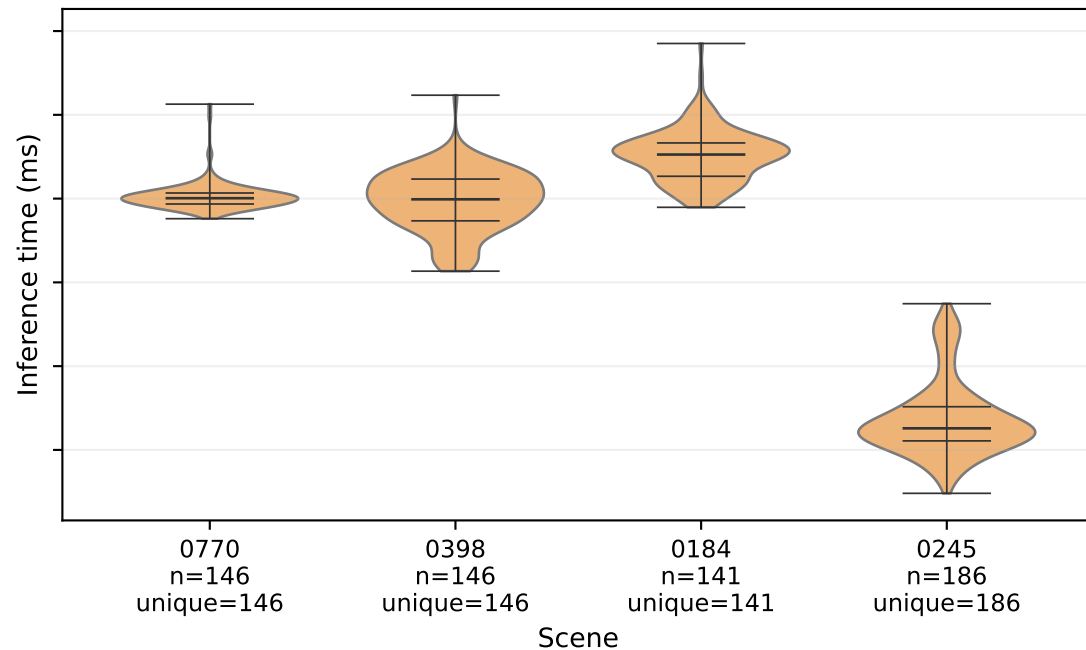
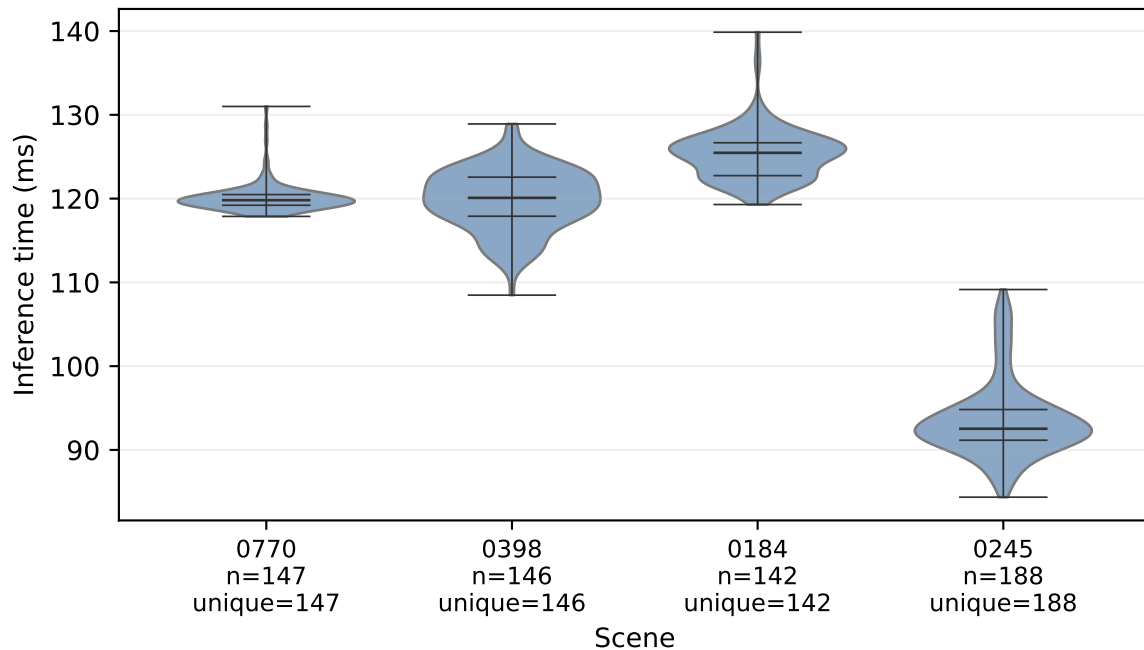


One execution per scene/mode; all completed non-warmup inferences. Bars: median, quartiles, min/max.
Shared y-scale across modes; frames are temporally dependent; P99 <1,000 samples is sparse.

centerpoint — single-model baseline

MPS off

MPS on

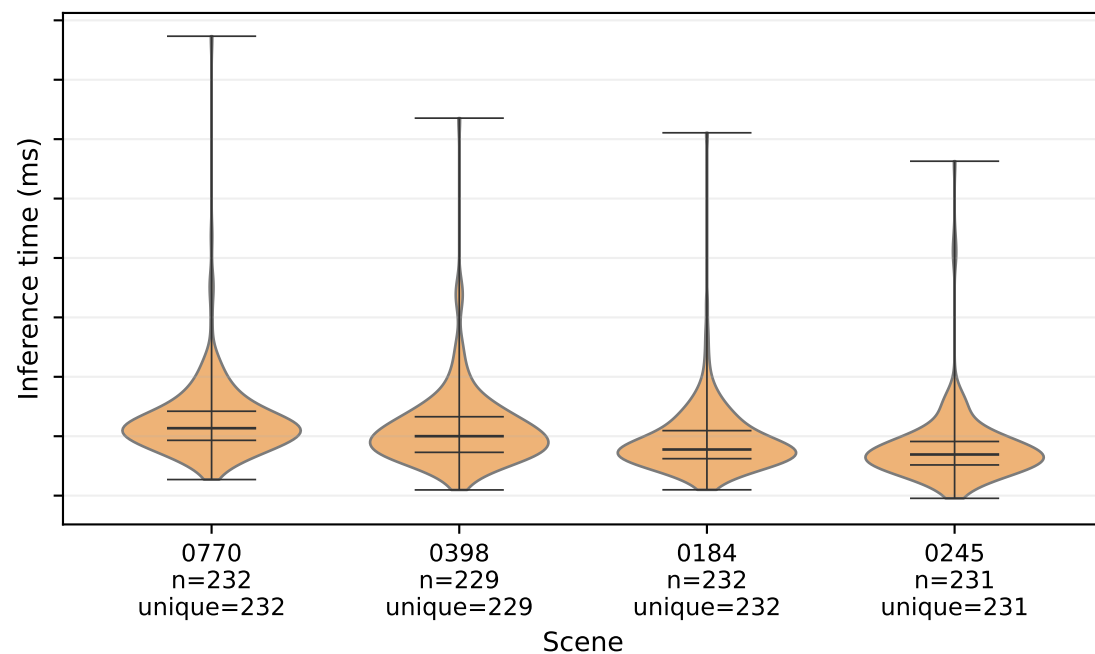
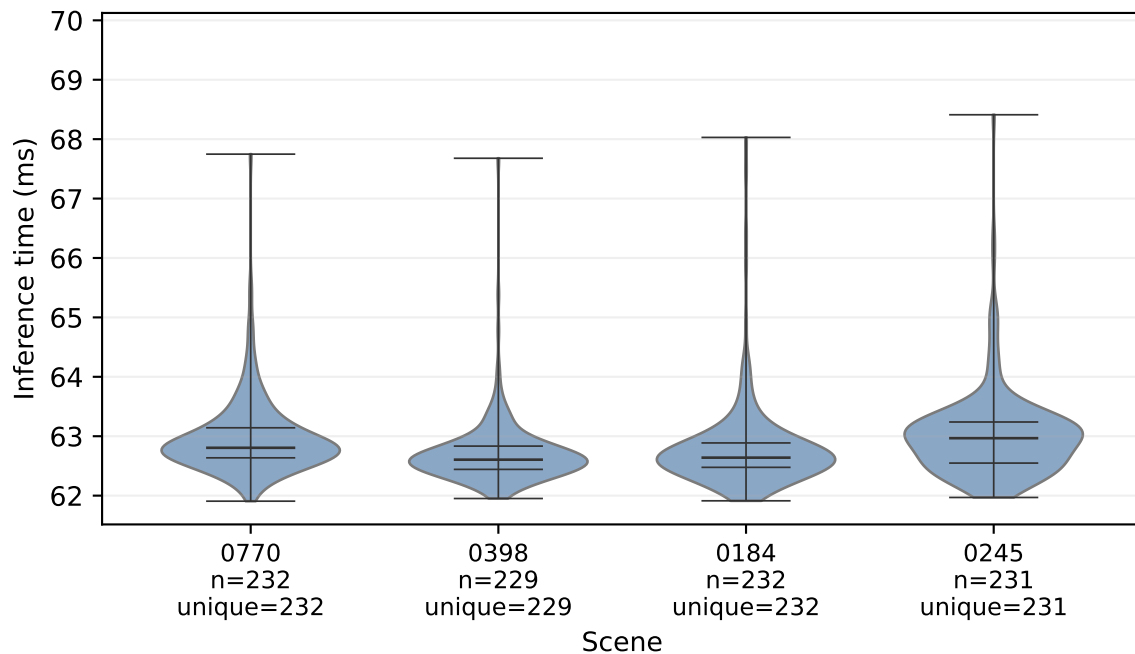


One execution per scene/mode; all completed non-warmup inferences. Bars: median, quartiles, min/max.
Shared y-scale across modes; frames are temporally dependent; P99 <1,000 samples is sparse.

dino — single-model baseline

MPS off

MPS on

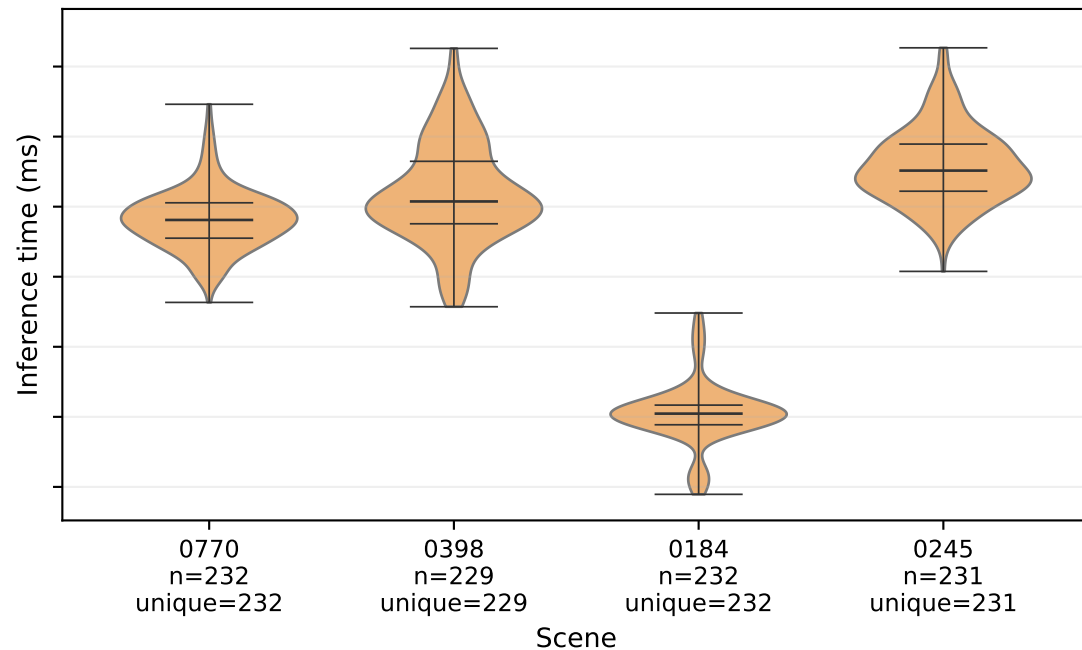
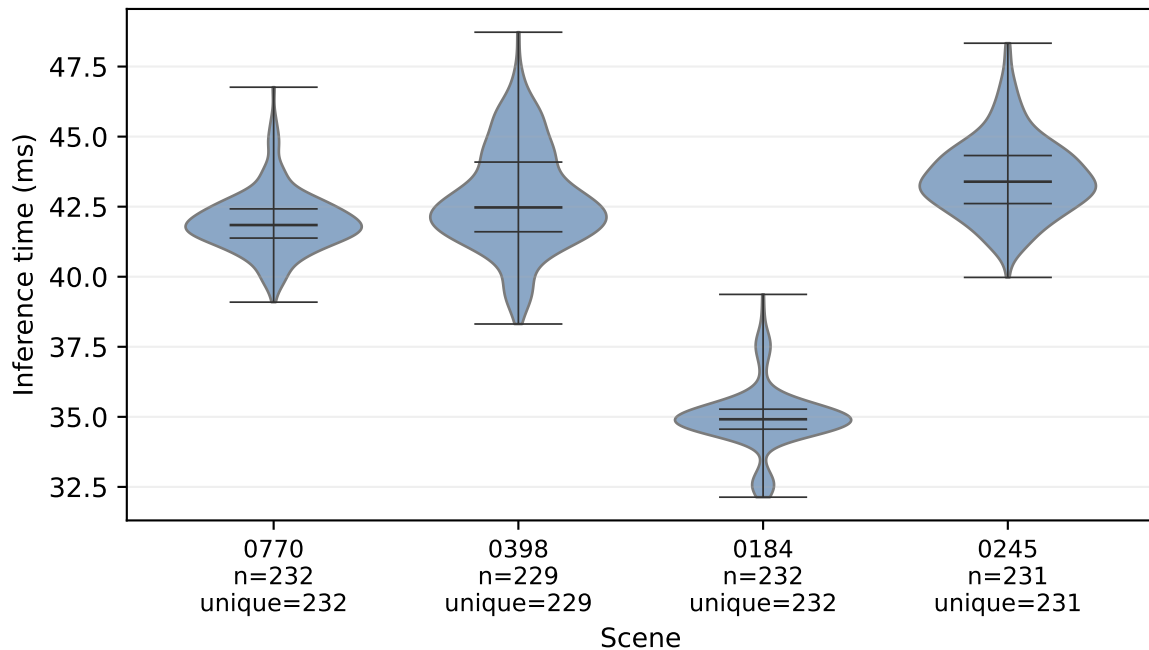


One execution per scene/mode; all completed non-warmup inferences. Bars: median, quartiles, min/max.
Shared y-scale across modes; frames are temporally dependent; P99 <1,000 samples is sparse.

mask-rcnn — single-model baseline

MPS off

MPS on

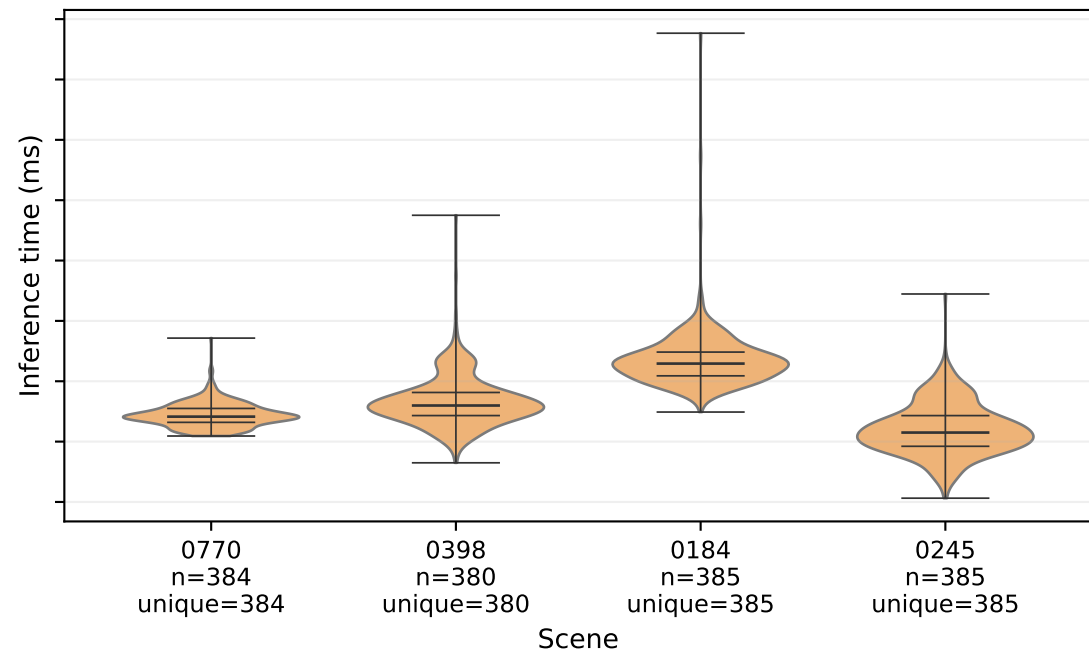
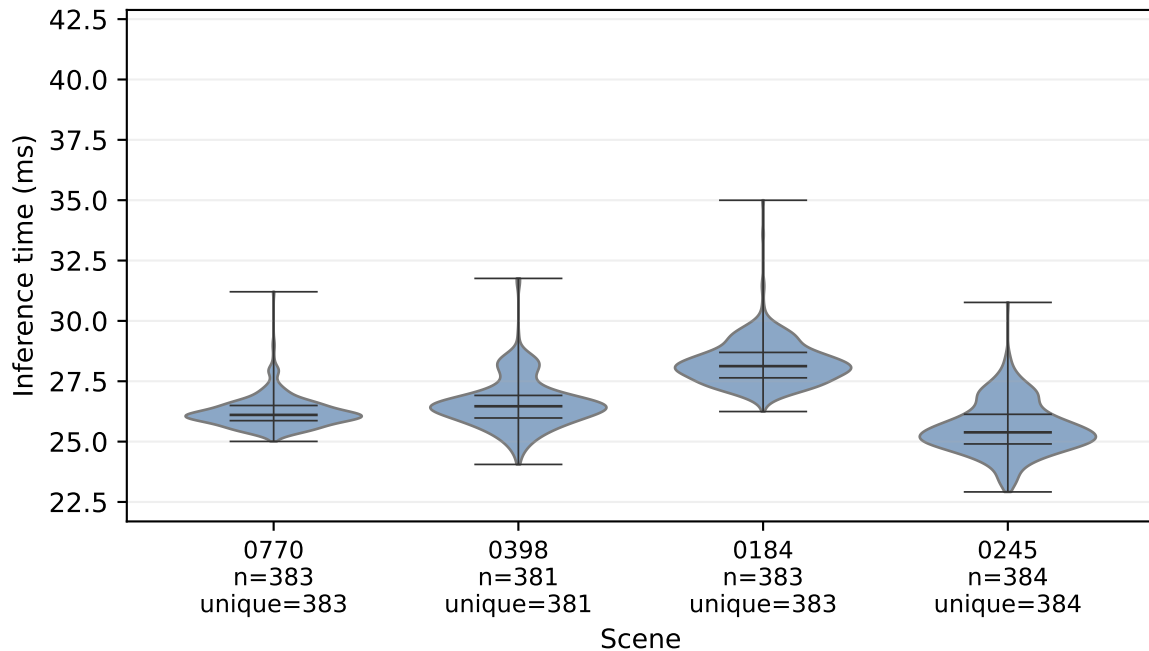


One execution per scene/mode; all completed non-warmup inferences. Bars: median, quartiles, min/max.
Shared y-scale across modes; frames are temporally dependent; P99 <1,000 samples is sparse.

pointpillars — single-model baseline

MPS off

MPS on

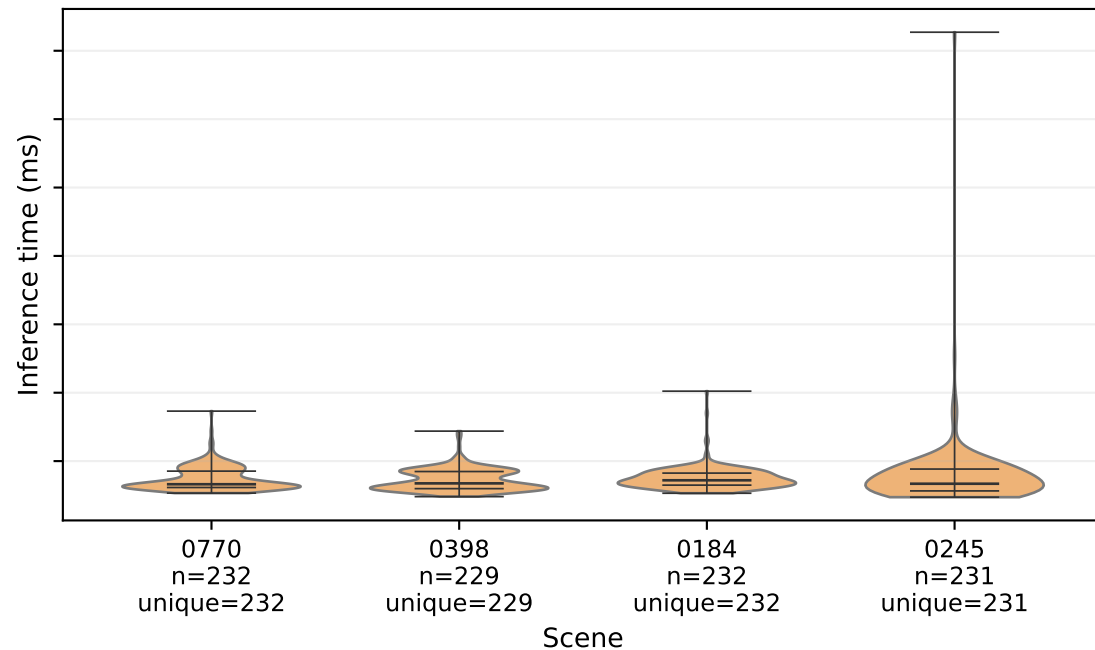
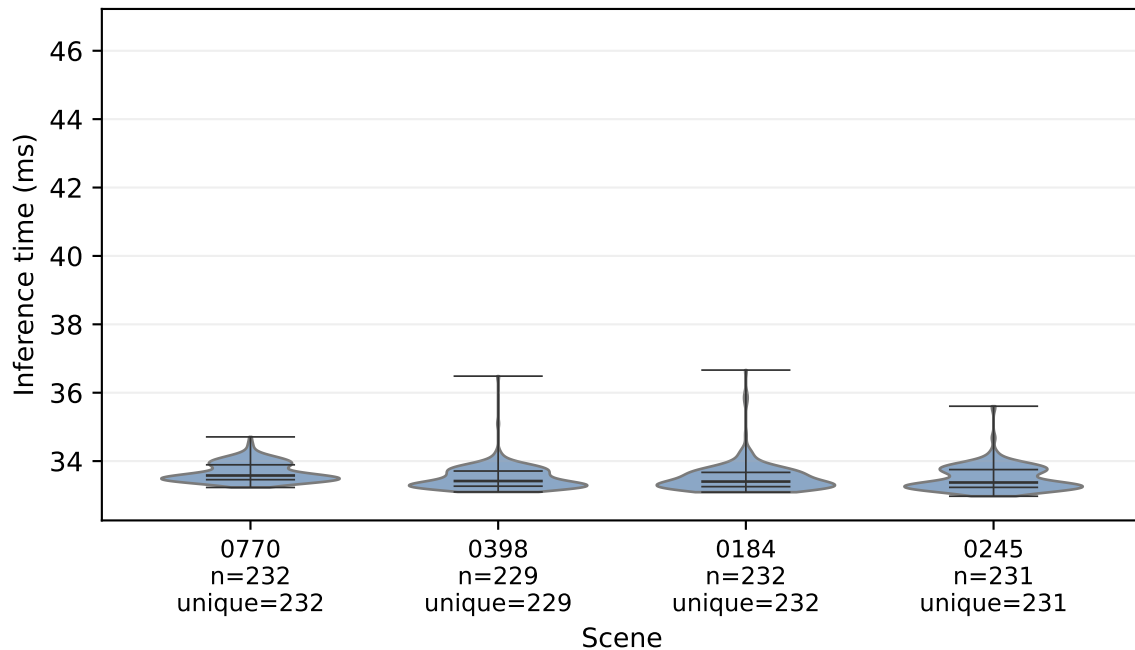


One execution per scene/mode; all completed non-warmup inferences. Bars: median, quartiles, min/max.
Shared y-scale across modes; frames are temporally dependent; P99 <1,000 samples is sparse.

vit-upernet — single-model baseline

MPS off

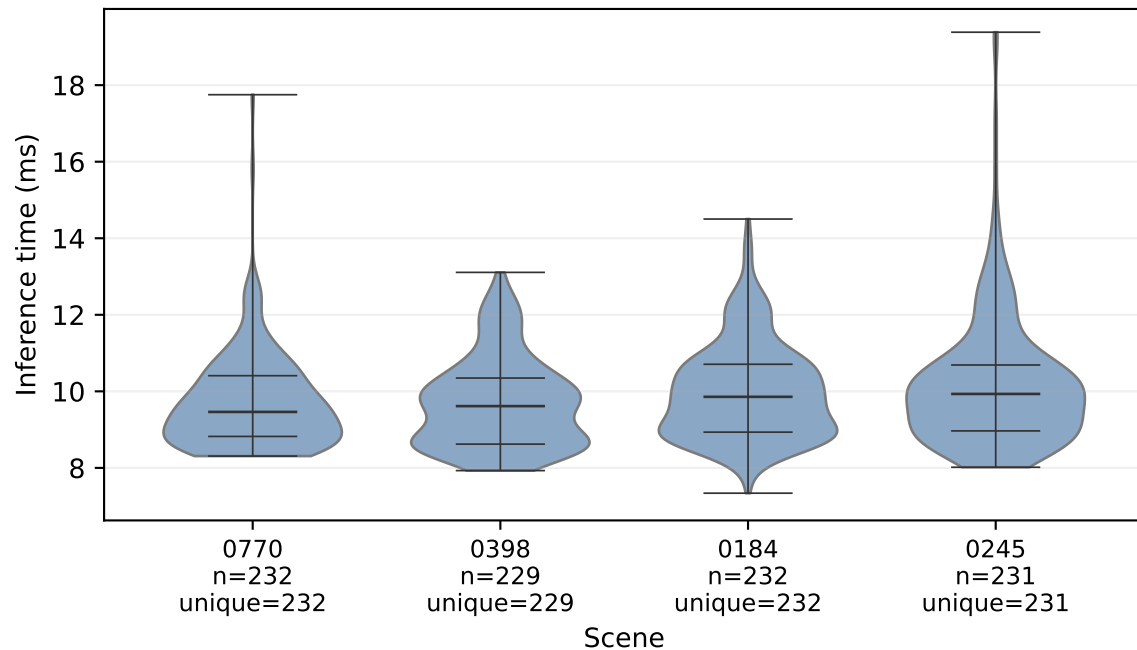
MPS on



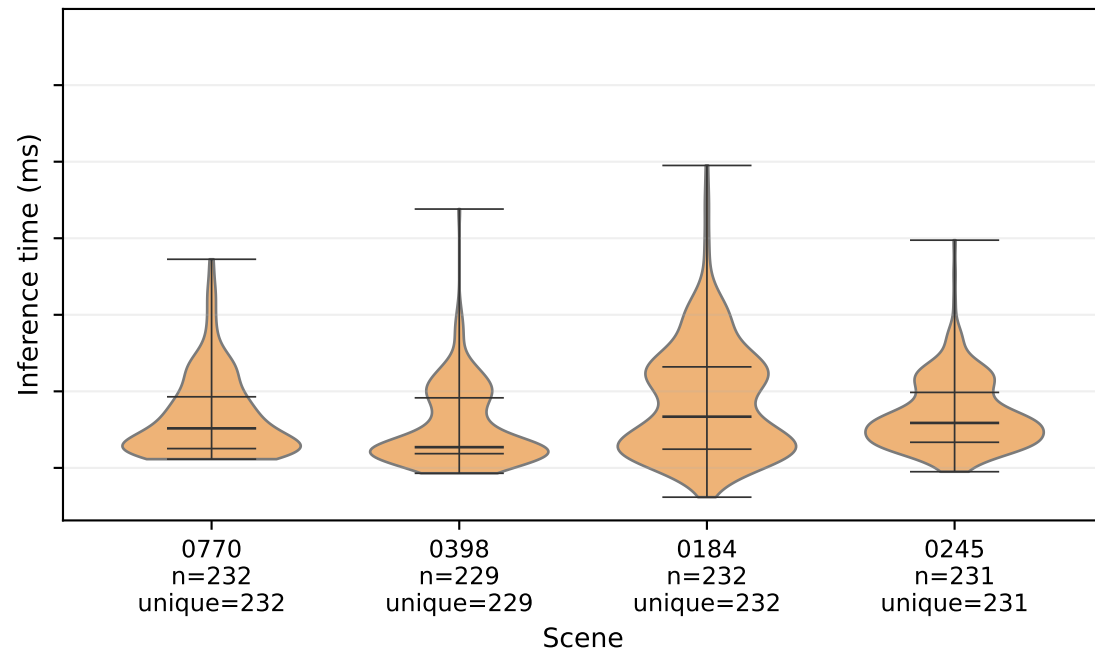
One execution per scene/mode; all completed non-warmup inferences. Bars: median, quartiles, min/max.
Shared y-scale across modes; frames are temporally dependent; P99 <1,000 samples is sparse.

yolov3 — single-model baseline

MPS off



MPS on



One execution per scene/mode; all completed non-warmup inferences. Bars: median, quartiles, min/max.
Shared y-scale across modes; frames are temporally dependent; P99 <1,000 samples is sparse.